

End Sem Exam 2020

Name \rightarrow Himanshu Gupta

Roll No \rightarrow BT17CSE093

Sem, Year \rightarrow 7th, Sem

Branch \rightarrow CSE

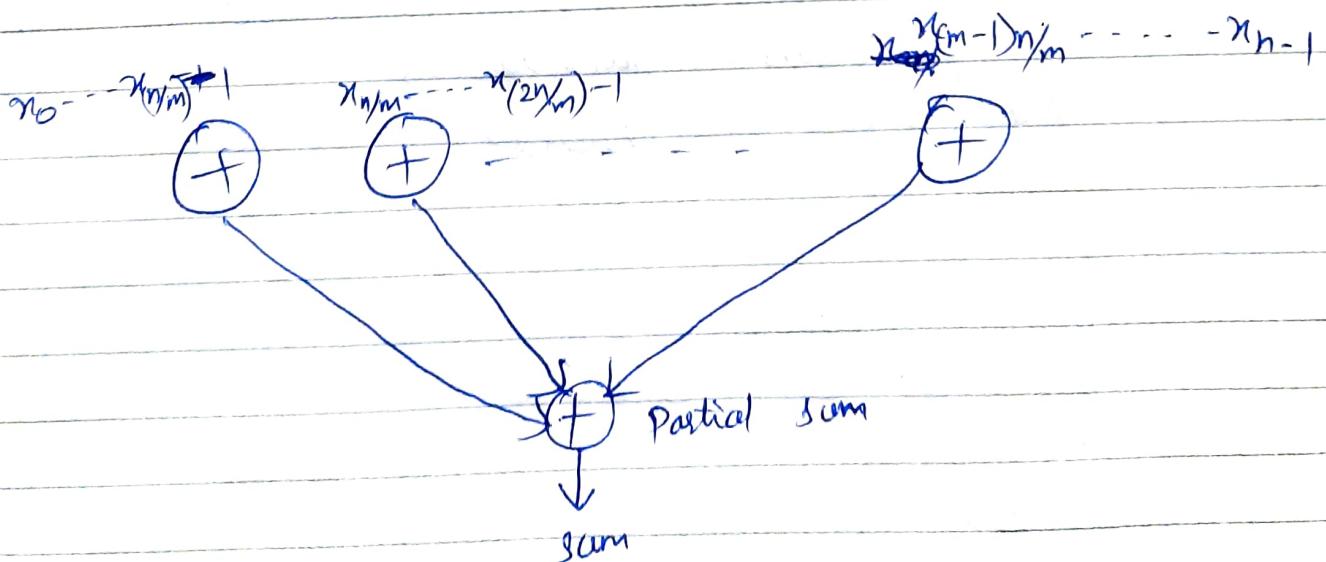
Course \rightarrow IPC

Date \rightarrow 26/11/2020

(a) Using Positioning strategy to solve the addition of numbers x_0, x_1, \dots, x_{n-1}

Steps -

- (1) Divide sequence into m parts of n/m numbers each and distribute ~~per~~ sequences to corresponding processors.
- (2) m processors can add each add one sequence independently to create partial sums
- (3) partial sums are added together on master to form final sum.



Sequential code

The code that uses separate `send()` and `recv()`.

Master:

```
s = n/m;           /* number of items on each slave */
for(i=0, n=0; i<m; i++) { n+=s;
    send(&numbers[n], s, Pj);
}
```

```
sum = 0;
for(i=0; i<m; i++) {   /* wait for results from slaves */
    recv(&part-sum, PANY);
    sum+= part-sum        /* accumulate partial sums */
}
```

3

Slave:

```
recv (number, s, Pmaster); // receive s items from master
part-sum = 0;
for(i=0; i<s; i++)      /* add numbers */
    part-sum += numbers[i];
send (&part-sum, Pmaster); /* send result to master */
```

4

Using Broadcast:

The code using Broadcast receiving:-

#

Master:

```
s = n/m;           /* no. of numbers for slaves */
broadcast (numbers, s, slave-group);
sum = 0;
```

```
for(i=0; i<m; i++) { /* wait for result from slaves */
    recv(&part_sum, PANY);
    sum+= part_sum;
}
```

3

Slave:-

```
broadcast(number, S, Pmaster); /* receive all numbers from master */
start = slave_number * S; // slave no. obtained earlier
end = start + S;
part_sum = 0;
for (i = start; i < end; i++)
    part_sum = part_sum + number[i]; // add numbers
send(&part_sum, Pmaster); // send sum to master.
```

So, in this sequential & parallel implementation is done using master & slave approach

$$\text{Speedup} = \frac{T_S}{T_P} = \frac{n-1}{n_m + m-2}$$

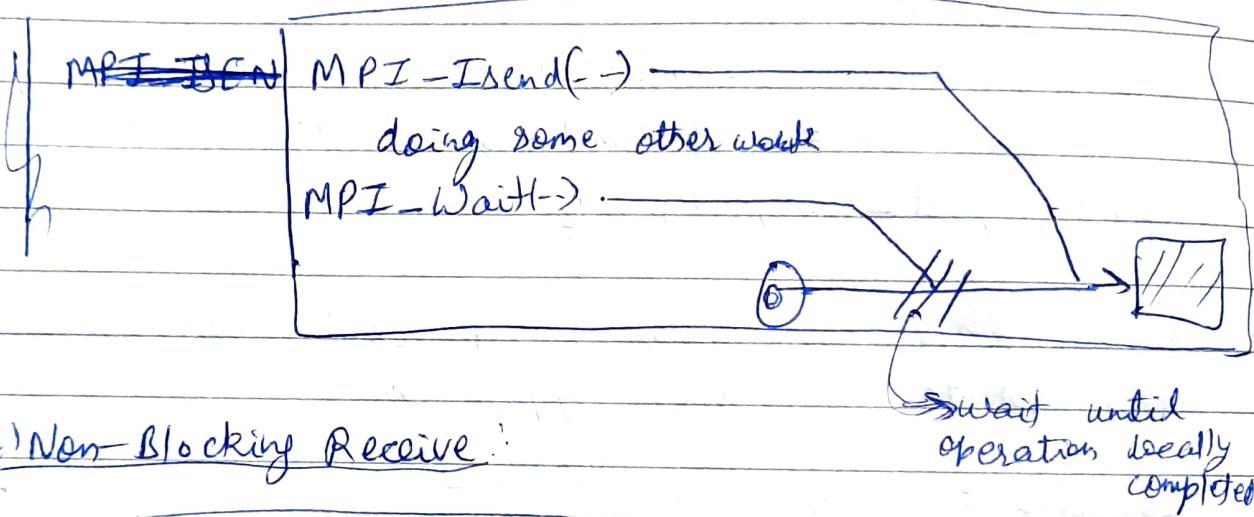
Ques 5) Blocking & Non-Blocking communication in MPI

Blocking communication is done using MPI-SEND() and MPI-RECV(). These functions do not return (i.e., they block), until the communication is finished. Simplifying somewhat, this means that the buffer passed to MPI-SEND() can be reused, either because MPI saved it somewhere or because it has been received by the destination. Similarly, MPI-RECV() returns when buffer has been filled with valid data.

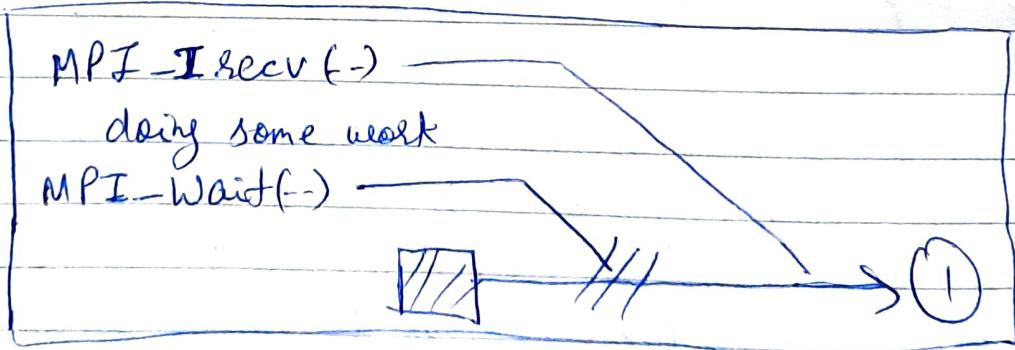
In contrast, non-blocking communication is done using `MPI_Isend()` and `MPI_Irecv()`. These function return immediate (i.e. they don't block) even if communication is not finished yet. You must call `MPI_Wait()` or `MPI_Test()` to see whether the communication has finished.

Example

(1) Non-Blocking Send:



(2) Non-Blocking Receive:



In C: other MPI functions

(i) `MPI_Isend (buf, count, datatype, dest, tag, comm)`
`OUT & request-handle);`
`MPI_Wait (INOUT & request-handle, & status);`

(ii) `MPI_Irecv(-)`

(iii) `MPI_Ibsend(- -)`

(iv) `MPI_Ibsend (-)`

→ Sync buffered.

(V) MPI-MAX()

(VI) MPI-MIN()

(VII) MPI-SUM()

(VIII) MPI-PROD()

(IX) MPI-LAND(-2)

(X) MPI-BAND(-2)

Q6) Assume size of A, B and C is 1000 elements.

(a) Thread blocks will generate = 4

(b) No. of warps in each block = 8

(c) Yes, there is control divergence, which is caused by the if statement in line 2.

if ($i < n$) $c - d[i] = A - d[i] + B - d[i/2]$

Since the total number of threads in the grid will be 1024, larger than the size of arrays, the last warp ~~will~~ will have divergence. First 8 thread in warp will take the true path and the other 24 threads will take false path.

(a) for ($i=0$; $i <= 100$; $i++$)
 for ($j=4$; $j <= 70$; $j++$)
 $H[i, j];$

$$(i) \text{ s: } A[i+2, j] = B[2i, j] - 3;$$

$$T: B[2i, j-1] = A[i, j+2] + 12;$$

for $s = (0, 0)$

$$\dots = \beta[6, 4] \quad \cancel{\dots}$$

$$B[6, 3] = \dots$$

so, there is anti-dependence of s on T

i.e. $\underbrace{s^a}_{} T$

$$\text{distance vector} = (0, 1)$$

$$\text{direction vector} = (0, 1)$$

$$\text{level} = 2.$$

for $s = (3, 6)$

$$A[5, 6] = \dots$$

$$= A[3, 8]$$

Hence, it is flow dependent

i.e. $\underbrace{s^f}_{} T$

Note,

$$\text{distance vector} = (2, -2)$$

$$\text{direction vector} = (1, -1)$$

$$\text{level} = 1$$

(ii) S: $A[i+2, j] =$

(ii) S: $A[i, j] = B[i+4, j-2] + B[i+2, j-3] + B[i, j+3];$

T: $B[i, j] = C[i, j] + A[i+1, j-2] + 2;$

for

$S(3, 4)$

S: $\dots = B[3, 7]$

T: $B[3, 4] = \dots$

do J

Here is $S \delta^a T$ (anti-dependence)

distance vector = $(0, 2)$

direction vector = $(0, 1)$

level = 2.

$T(3, 5) \delta^a S(4, 4)$ for S: $A[4, 4] = \dots - A[5, 3]$

distance vector = $(1, -1)$

direction = $(1, -1)$

level = 1

$S(3, 6) \delta^a T(7, 4)$ S: $\dots = B[7, 4]$
T: $-B[3, 6]$

distance vector = $(4, -2)$

dissection = $(1, -1)$

level = 1

$S(3, 7) \delta^a T(5, 4)$ S: $\dots = B[i+2]$

T: $-B[i, j] = \dots$

distance vector = $(2, -3)$

dissection = $(1, -1)$

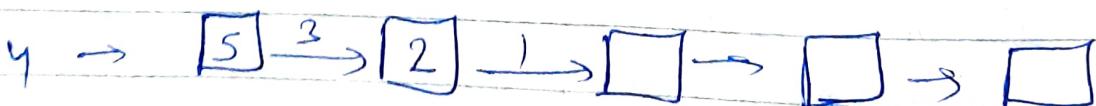
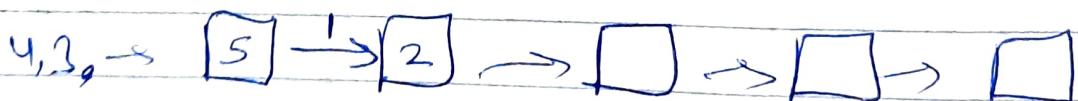
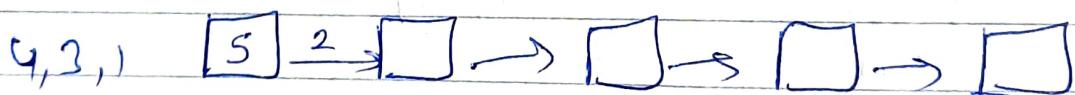
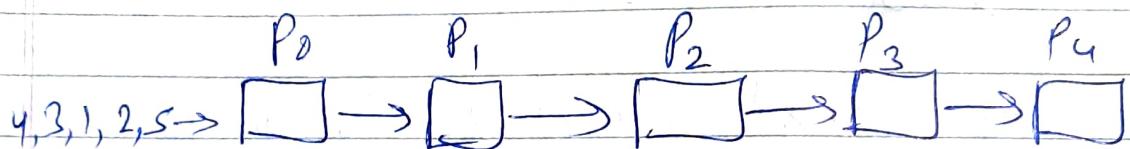
level = -1

Ans 4) We need to sort ~~the elements of array.~~
we can use Insertion sort with bits line.

Algorithm:-

- ① Insert data elements to one of the processor
- ② Each processor on sorting element will compare itself with given number.
- ③ If given number is large, it will keep the number and pass smaller number
- ④ else, pass the given number.

Example : array = [4, 3, 1, 2, 5] & no. of processors = 5



Assuming compare & exchange as single computational step

$$t_s = (n-1) + (n-2) + \dots + (2) + 1 = \frac{n(n+1)}{2}$$

whereas, parallel Algorithm has

$$n + n - 1 = 2n - 1 \quad (\text{pipeline cycles})$$

$$t_{\text{computation}} = 1$$

$$t_{\text{communication}} = 2(t_{\text{copy}} + t_{\text{data}})$$

&

$$\text{So, } t_p = O(n)$$

and serial, $T_s = n^2$

$$\boxed{\text{speedup} = \frac{T_s}{T_p} = \frac{n^2}{n} = n}$$

No. of Assignment = 4.950

② ~~RoundRobin~~ (static, 1)

thread 0 : 0 4 8

thread 1 : 1 5 9

thread 2 : 2 6 10

thread 3 : 3 7 11

issue:- Not utilising the cache well.
0, 1, 2, 3 will be stored nearly in sequential form in storage memory (assuming numbers)

(1) Schedule Static (Static, 10)

thread 0: [0-9] [40-49] [80-89]

thread 1: [10-19] [50-59] [90-99]

thread 2: [20-29] [60-69] [100-109]

thread 3: [30-39] [70-79] [110-119]

Issue: Does not take care of the dynamic nature (i.e. which computation takes more time, which takes less, allocating sets accordingly).

(2) Schedule (dynamic, 10)

thread 0: [0-9] [40-49] [60-69]

thread 1: [10-19] [50-59] [70-79] [100-109]

thread 2: [20-29] [60-69] [90-99]

thread 3: [30-39] [60-69] [80-89]

→ More balanced Allotment.

(u) Schedule (dynamic, 20)

thread 0 :- [0-9]

1 : [20-29]

2 : [40-59]

3

4 : [60-79]

now it

chunk size ↑

static size ↑

dynamic nature ↓